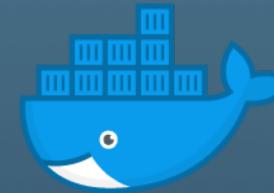


# DBA MASTERY



Exploring SQL Server  
containers on Docker and  
Kubernetes





# Carlos Robles

Principal Consultant, DBA Mastery



/croblesdba



@dbamastery



crobles@dbamastery.com

## Experience

Microsoft Data Platform MVP  
Over 10 years of experience  
Multi platform DBA

## Community

International speaker, author, blogger, mentor  
Guatemala SQL Server community leader  
Simple Talk, SQL Server Central and MSSQL Tips  
author

## DBA Mastery

SQL Server tips, scripts, best practices and more



MAXDOP Calculator

Azure Data Studio wait stats widget

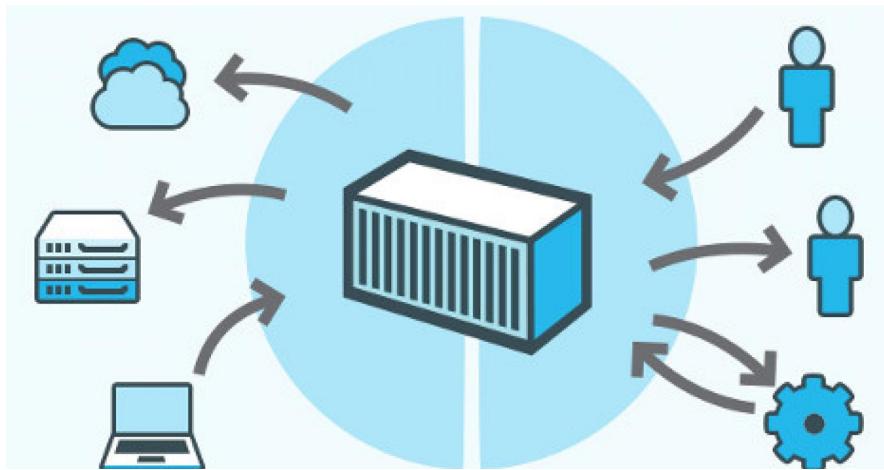
PerfMon for DBA's

MSDB tuning

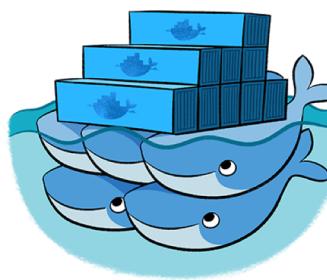
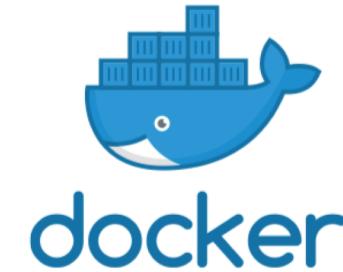
# AGENDA

- Introduction to Docker
  - Images
  - Containers
  - Docker architecture
  - VM's vs Containers
  - Advantages
- The SQL Server docker image
- The SQL Server Dockerfile
- Running a SQL container
  - Managing containers
  - Demo (Linux \ Unix + Windows)
- Introduction Kubernetes
  - Kubernetes architecture
  - Advantages
  - Demo





Containers are the future!



# DOCKER



- From Docker docs:

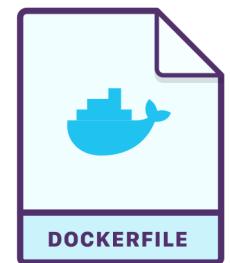
*Docker is an open platform for developing, shipping, and running applications. Docker enables you to separate your applications from your infrastructure so you can deliver software quickly.*

*With Docker, you can manage your infrastructure in the same ways you manage your applications.*



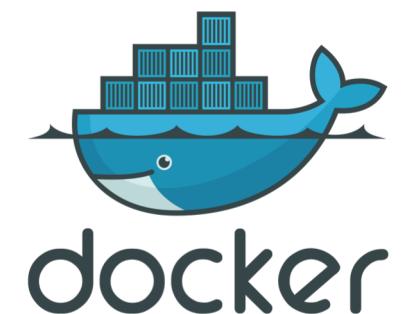
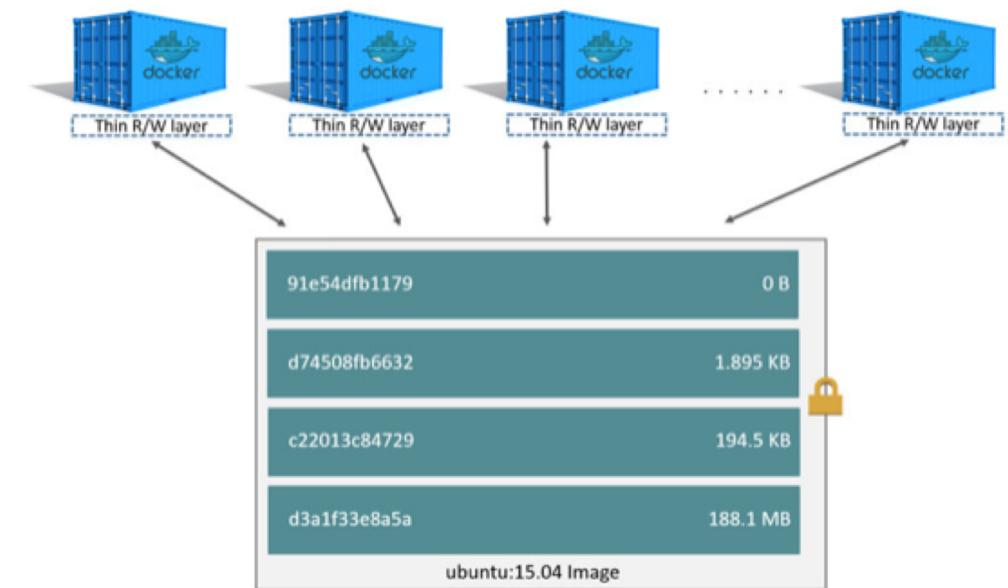
# IMAGES

- Is a read-only template with instructions for creating a Docker container
- Images are created using a Dockerfile
- A snapshot of a set of files required to run an application
- Portable and consistent
- A new image can be created from an existing image (make your own)
  - SQL Server for example, based on Ubuntu or RedHat

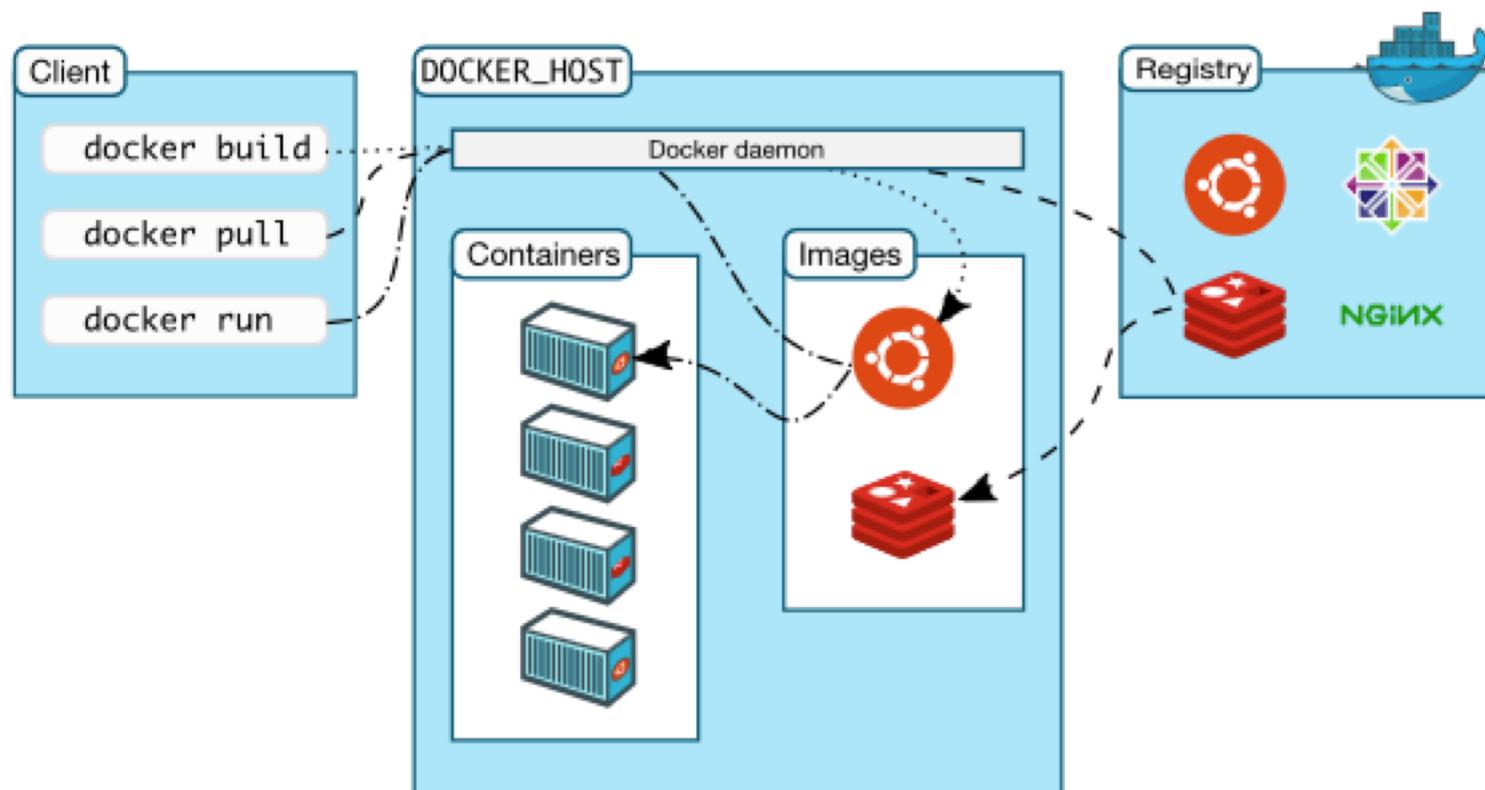


# CONTAINERS

- The runnable instance of a Docker image
- A standardized unit of software
- Container is nothing more than a program
- Writable layer and shared read-only layer
  - Small storage footprint
- Containers has full access to all resources
- Volumes = Persistent storage



# DOCKER ARCHITECTURE



# VM'S VS CONTAINERS

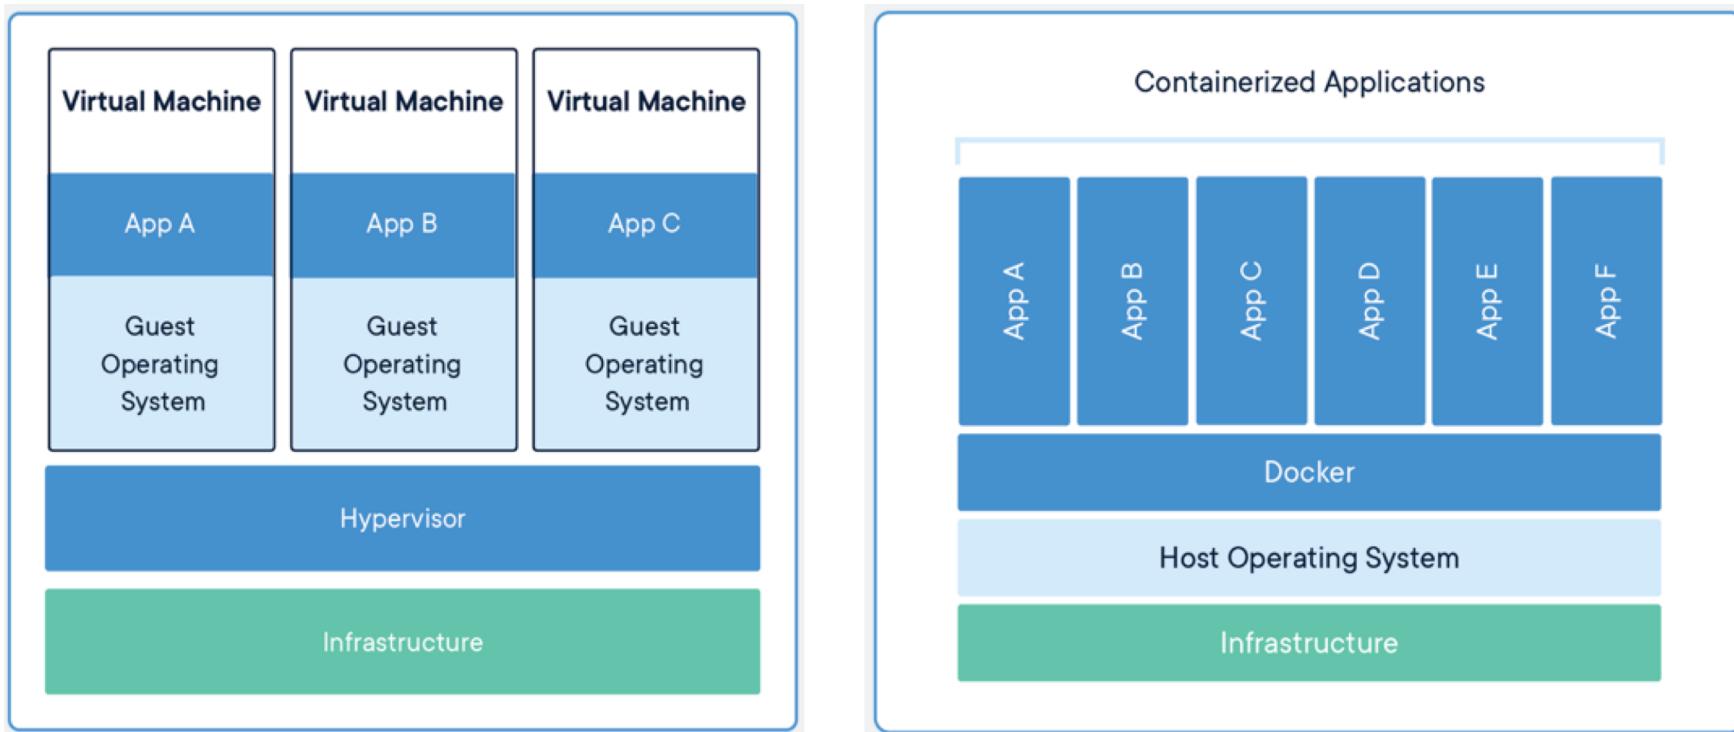


- 
- Virtualization +15 years
  - Sometimes heavyweight
  - Hardware virtualization
  - Each VM has an entire OS



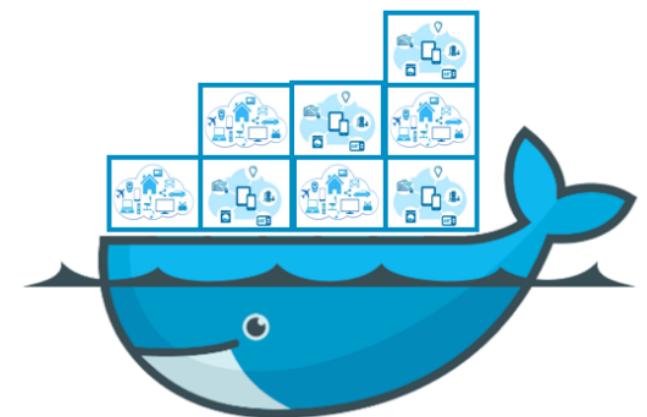
- 
- No installation
  - Lightweight
  - OS virtualization
  - All containers run in the same host OS





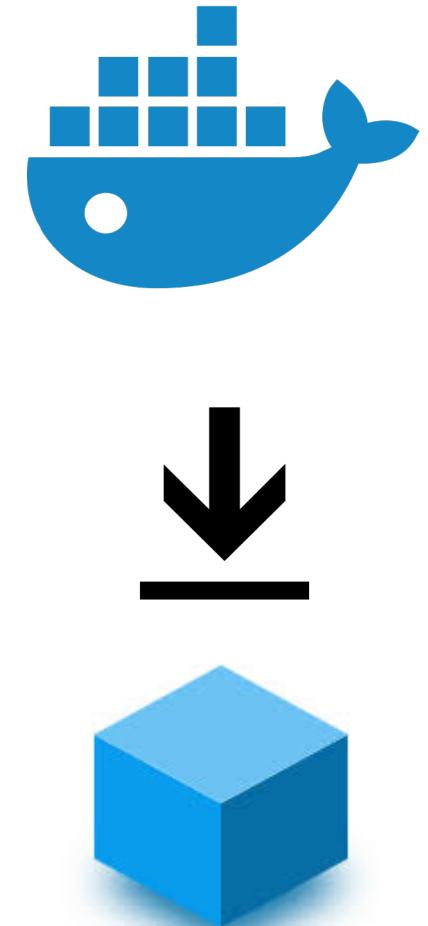
# ADVANTAGES

- Easy to use
- Agile application creation and deployment
- CI\CD – DevOps
- Resource isolation and better utilization
- Quick start \ stop time
- Cloud and OS portability
- Environmental consistency across all environments



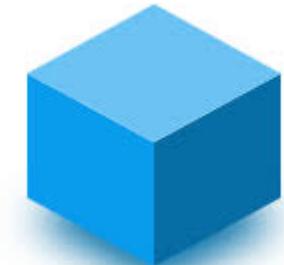
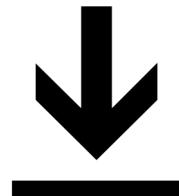
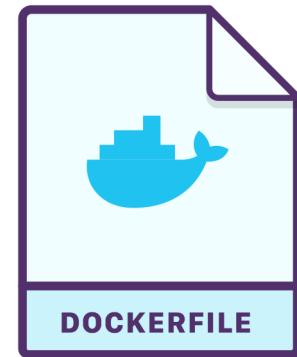
# SQL SERVER IMAGE

- [Docker Hub](#) – [Microsoft container registry](#)
- SQL Server 2017
  - Just Ubuntu from RTM to latest CU and GDR
- SQL Server 2019 (CTP)
  - Ubuntu and RedHat
  - From CTP 1.0 to latest
- SQL Server is pre-installed (standard)
- Backups are compatible between all platforms



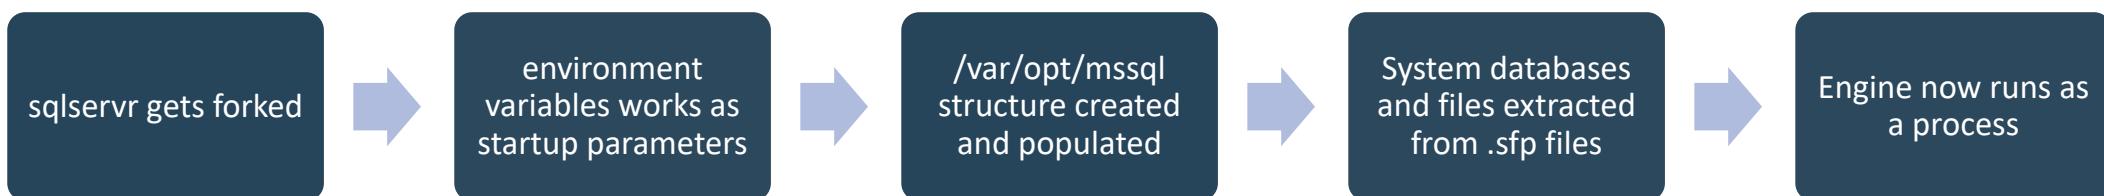
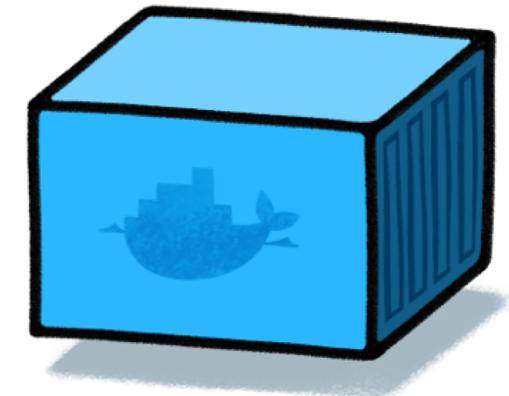
# SQL SERVER DOCKERFILE

```
FROM ubuntu:16.04  
  
EXPOSE 1433  
  
COPY ./install /  
  
CMD ["/opt/mssql/bin/sqlservr"]
```

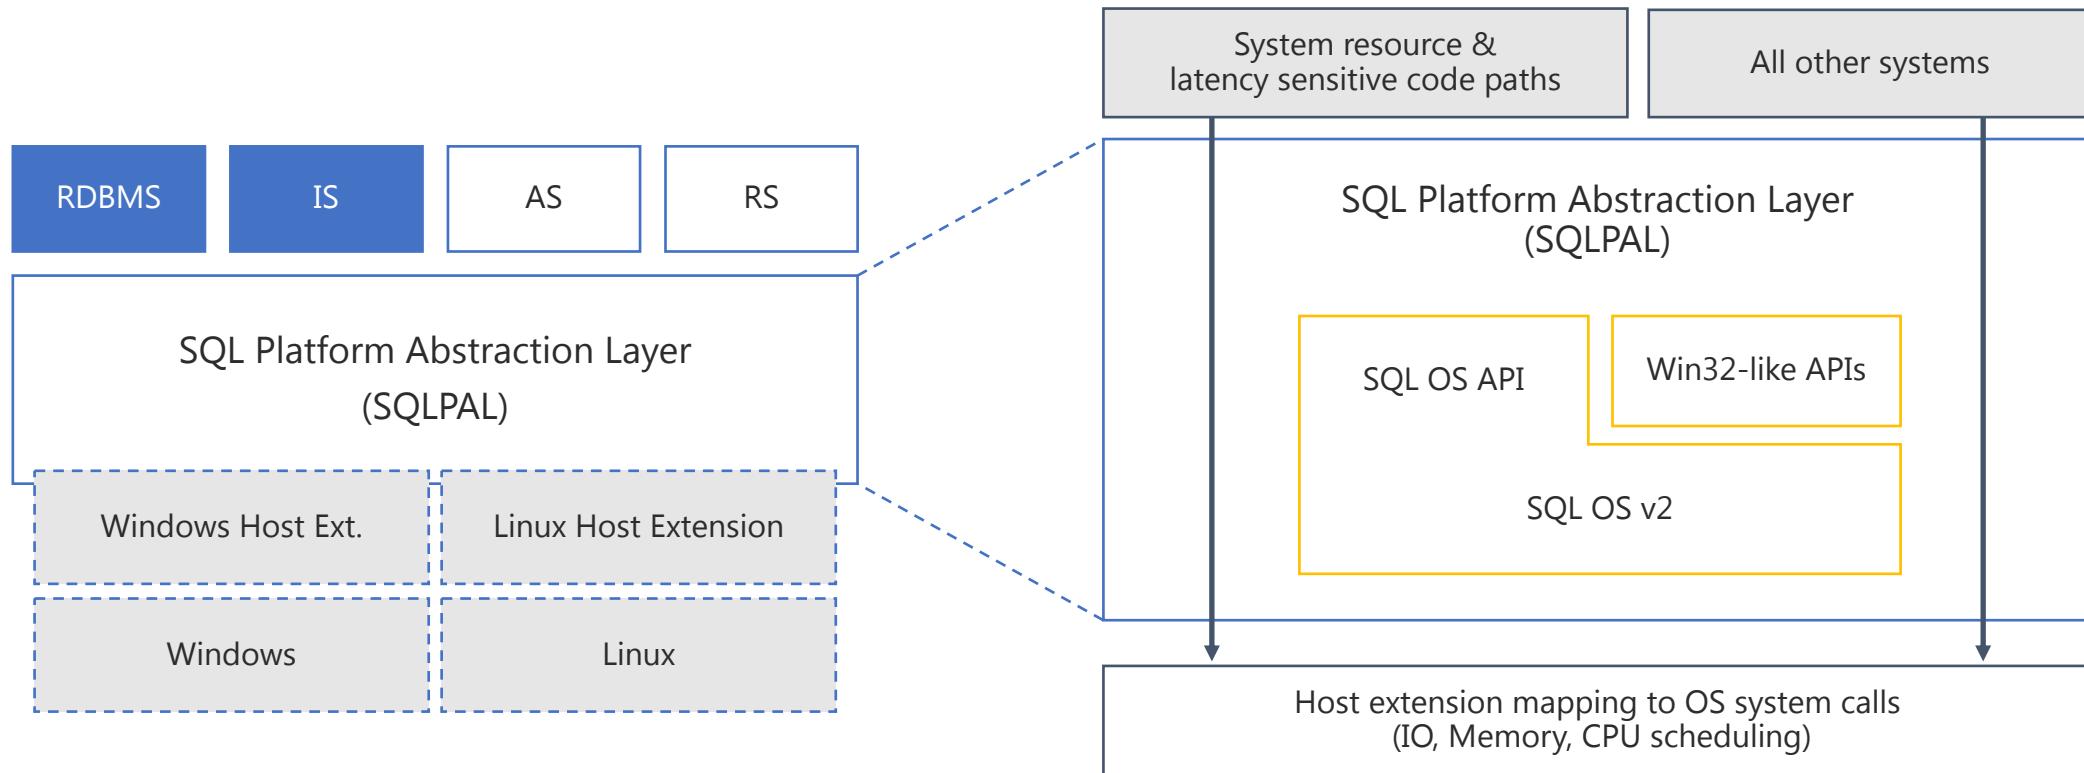


# RUNNING A SQL CONTAINER

```
docker run \
--name TEC_GT \
--env 'ACCEPT_EULA=Y' \
--env 'MSSQL_SA_PASSWORD=T3cG7-R0cks' \
--publish 1400:1433 \
--detach mcr.microsoft.com/mssql/server:2017-CU15-ubuntu
```



# SQL SERVER LINUX ARCHITECTURE



# SQL SERVER 2019 ON LINUX

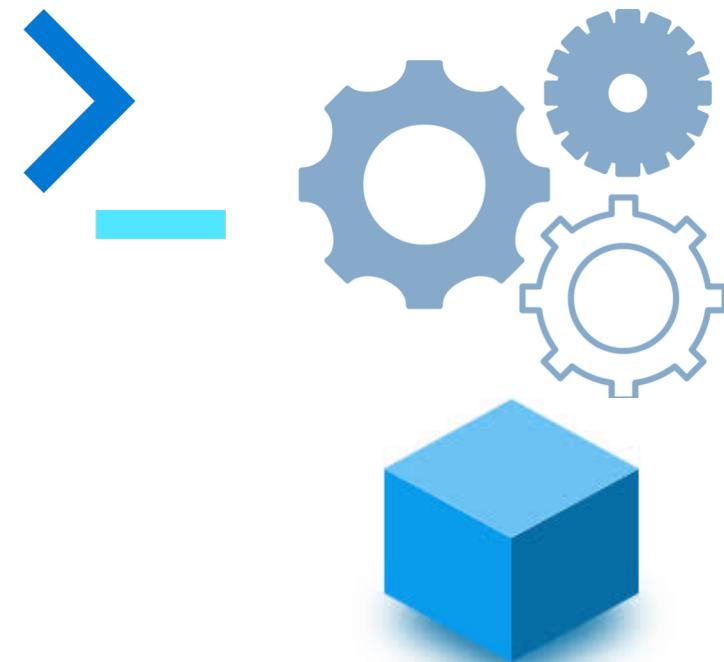
		Windows	Linux
<b>Editions</b>	Developer, Express, Web, Standard, Enterprise	●	●
<b>Services</b>	<b>Database Engine, Integration Services, SQL Server Agent</b>	●	●
	Analysis Services, Reporting Services, MDS, DQS	●	
	Maximum number of cores	Unlimited	Unlimited
	Maximum memory utilized per instance	24 TB	12 TB
	Maximum database size	524 PB	524 PB
<b>Mission critical performance and HADR</b>	Basic OLTP (Basic In-Memory OLTP, Basic operational analytics)	●	●
	Advanced OLTP (Advanced In-Memory OLTP, Advanced operational analytics, adaptive query processing)	●	●
	SQL Server Replication	●	● NEW
	Basic high availability (2-node single database failover, non-readable secondary)	●	●
	Advanced HA (Always On - multi-node, multi-db failover, readable secondaries)	●	●
<b>Security</b>	Basic security (Basic auditing, Row-level security, Data masking, Always Encrypted, <b>Active Directory Authentication</b> )	●	●
	Advanced security (Transparent Data Encryption)	●	●
	PolyBase	●	● NEW
<b>Data warehousing</b>	Basic data warehousing/data marts (Basic In-Memory ColumnStore, Partitioning, Compression)	●	●
	Advanced data warehousing (Advanced In-Memory ColumnStore)	●	●
	Advanced data integration (Fuzzy grouping and look ups)	●	●
<b>Tools</b>	Windows ecosystem: Full-fidelity Management & Dev Tool (SSMS & SSDT), command line tools	●	●
	Linux/OSX/Windows ecosystem: Dev tools (VS Code), DB Admin GUI tool, command line tools	●	●
<b>Developer</b>	Programmability (T-SQL, CLR, Data Types, JSON, Graph)	●	●
	Distributed Transactions	●	● NEW
	Machine Learning Services	●	● NEW

# DEMO



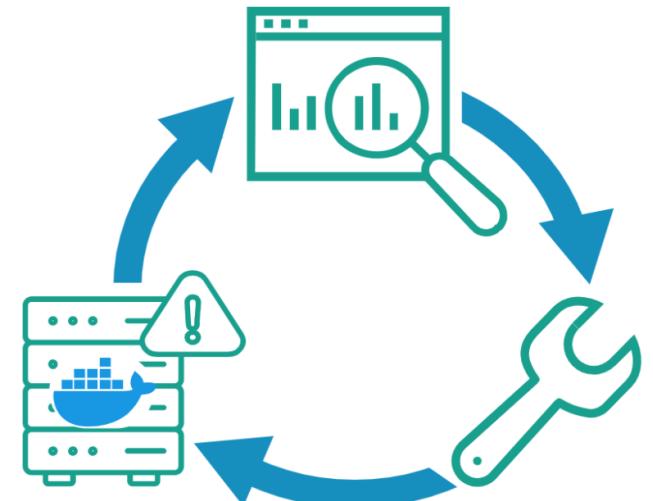
# MANAGING CONTAINERS

- docker pull
- docker run
- docker start | stop
- docker image | container
- docker rm | rmi
- docker exec
- docker build
- docker logs
- docker inspect
- docker volume
- docker save



# USE CASES

- Local development
- Troubleshooting
- Demonstrations
- Eliminates shared environments
- Eliminates resource contention
- Temporal environments
  - No installation \ patching



# KUBERNETES



kubernetes

- From Kubernetes docs:

*Kubernetes is a portable, extensible open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation.*

*It also make possible the container orchestration for automating application deployment, scaling, and management.*



# KUBERNETES ARCHITECTURE

- Masters
  - Multiple moving parts \ processes
  - Runs on a single node in the cluster
  - Tells others what to do – desired state
- Nodes
  - Do the work, runs applications
  - Aka “minions”
  - Reports the state back up to the master



- Pods
  - Containers runs inside of pods
  - Can have one or more pods within a node
- Services
  - Hiding multiple pods behind a service IP address
- Deployments
  - Declarative model
    - Desired state (number of POD's)
  - Manifest file (YAML, JSON)



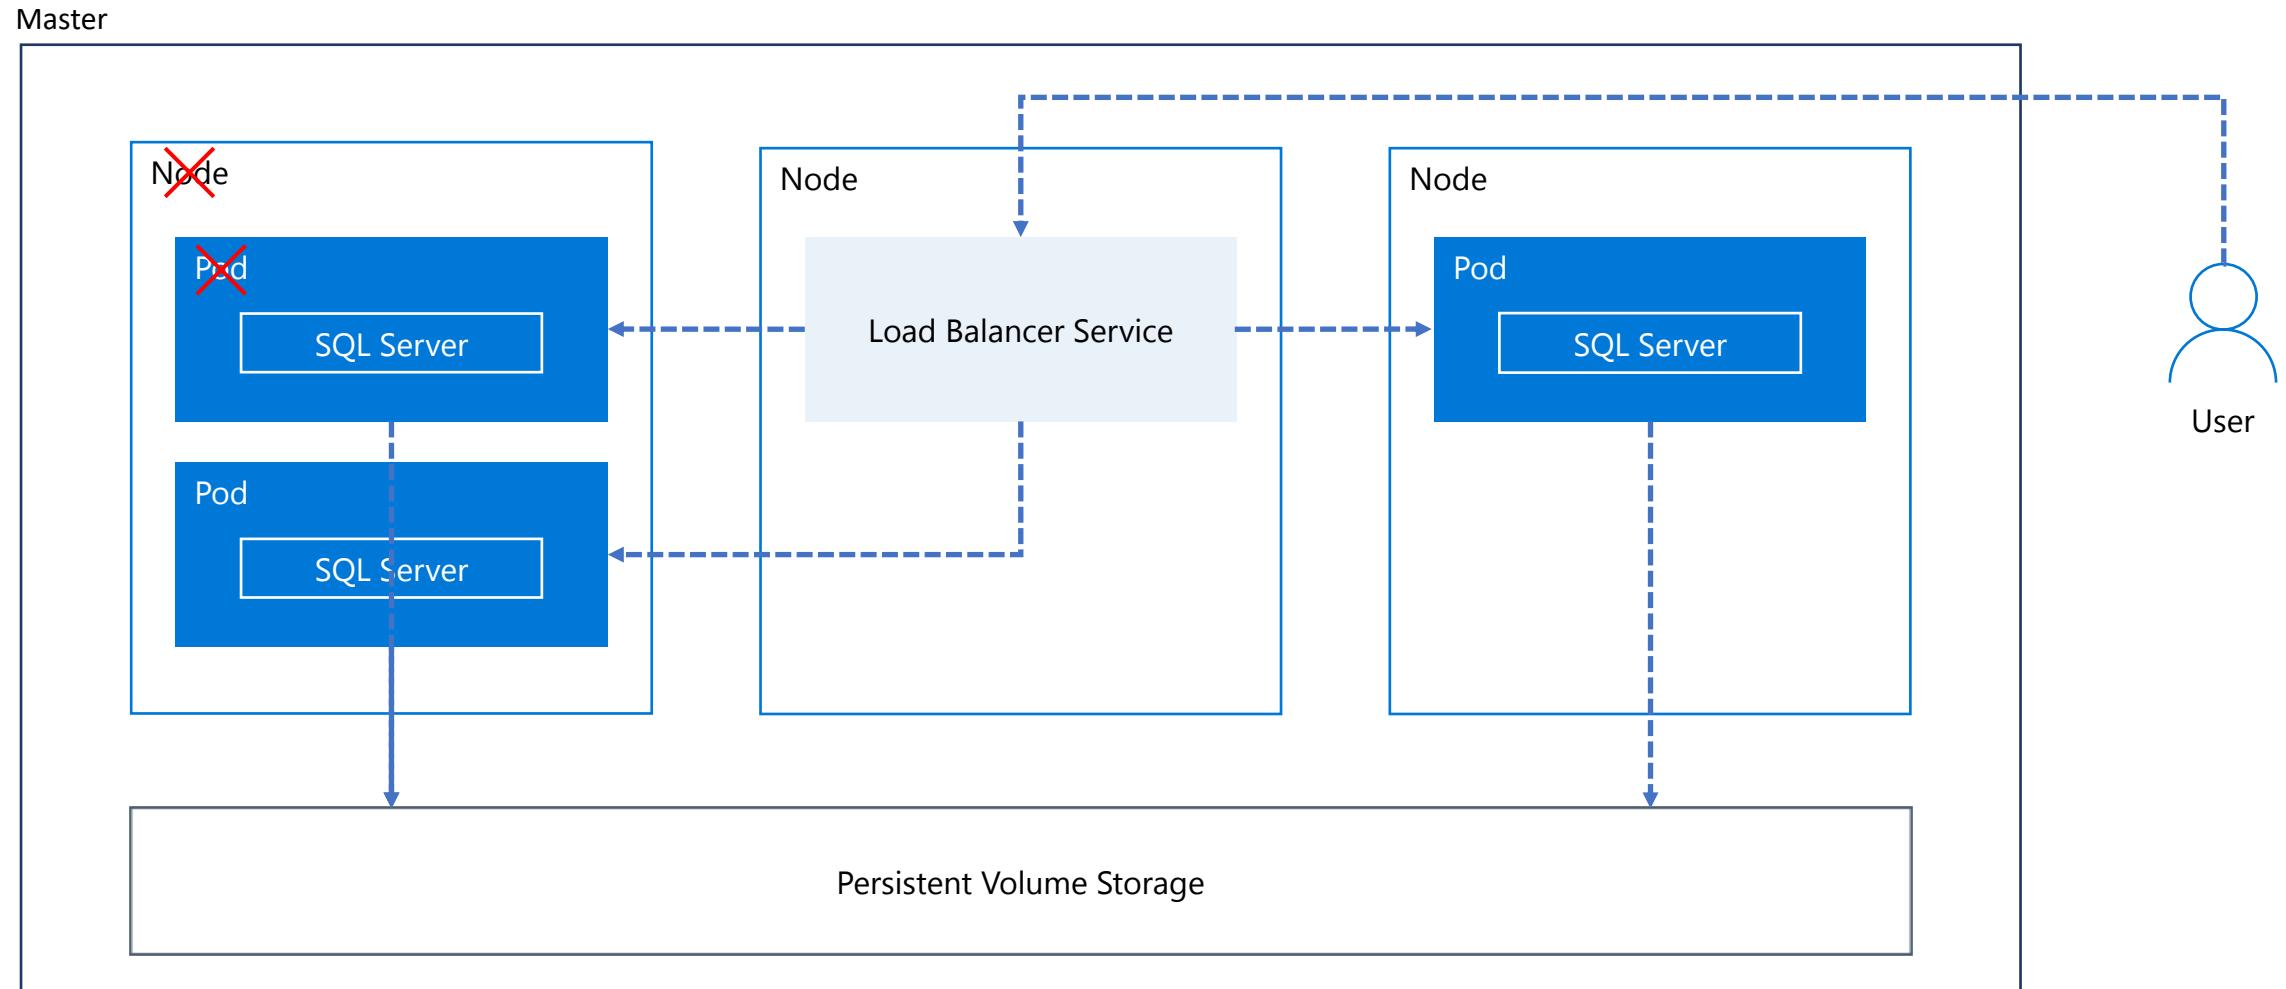
# ADVANTAGES

- Easy to use – Declarative configuration
- Self healing – Built in HA
- Auto scale
- Platform agnostic
- Compute and storage layer are separate
- Load balancing
- CI\CD – DevOps workflow



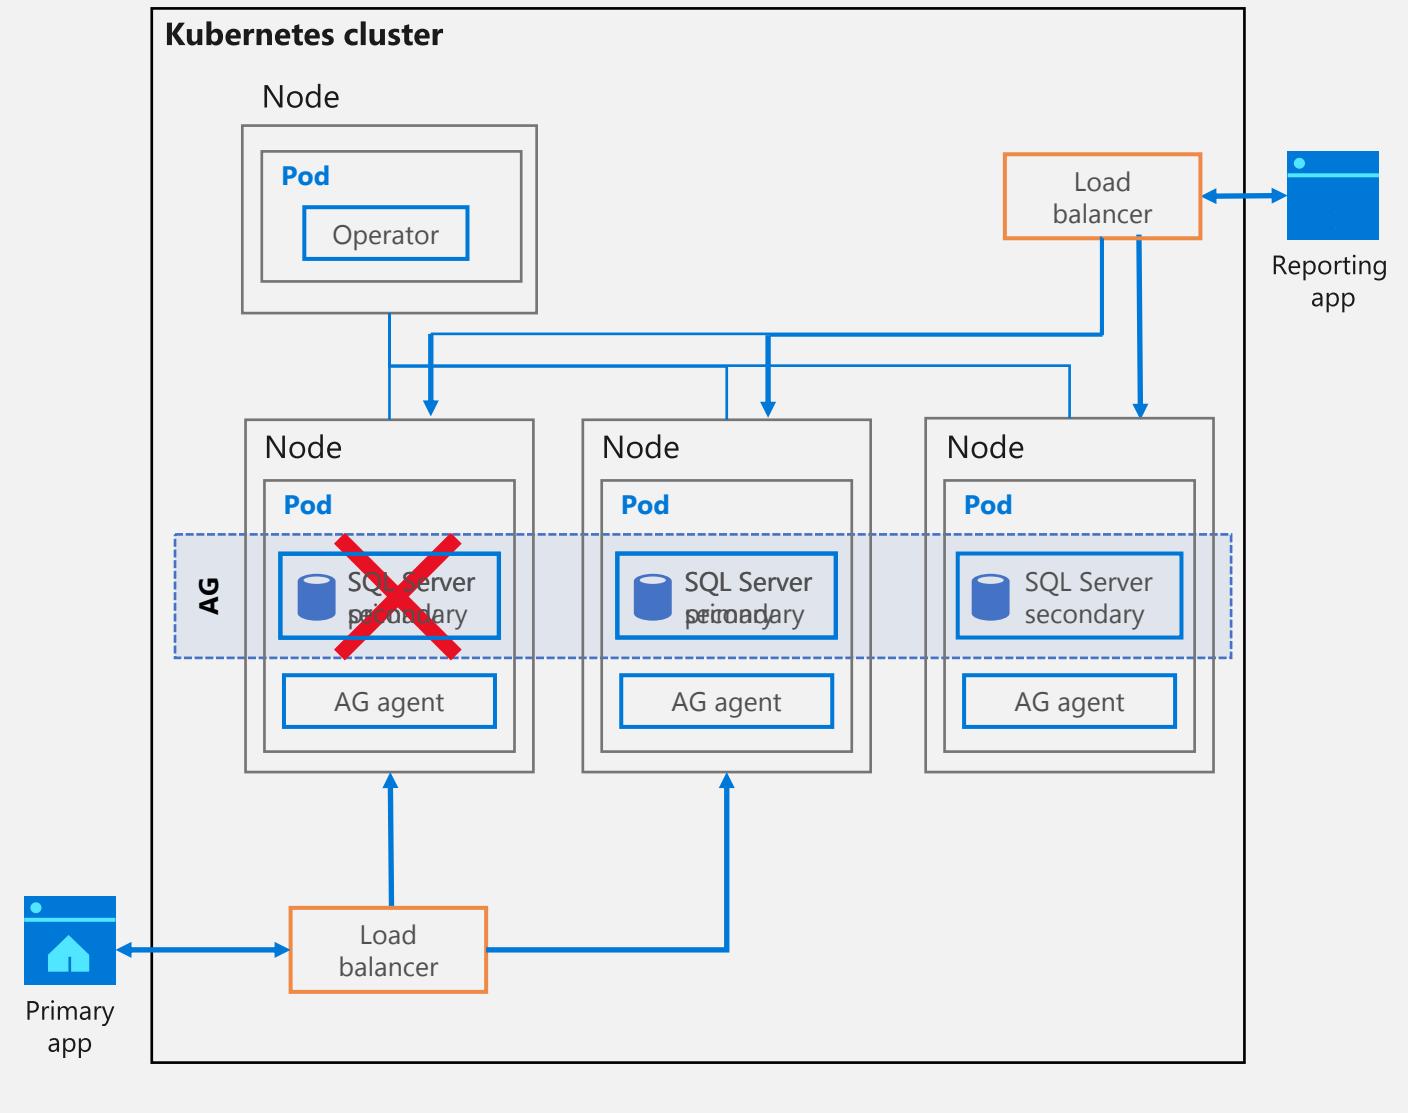
# DEMO





## Availability groups on Kubernetes

- SQL Server/k8s failover integration
- Operator deployment
- AG concepts all apply
- Load Balancer for Primary App
- Load Balancer for Secondary Replica Readers

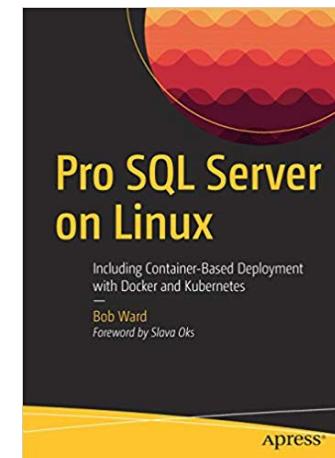


# QUESTIONS ?



# REFERENCES

- **Official documentation**
  - [Docker Docs](#)
  - [Kubernetes Docs](#)
- **SQL Server**
  - [SQL Server workshops](#)
  - [SQL Server samples](#)
- **Books**
  - [Docker Deep dive](#)
  - [The Docker book](#)
  - [Kubernetes: Up and Running](#)
  - [The Kubernetes book](#)
  - [Pro SQL Server On Linux by Bob Ward](#)
- **Pluralsight Courses**
  - [Getting Started with Docker](#)
  - [Docker Deep Dive](#)
  - [Docker and Kubernetes: The big picture](#)
  - [Kubernetes Installation and Configuration fundamentals](#)
- **Microsoft Learning Courses**
  - [Kubernetes Learning Path](#)
  - [SQL Workshops](#)
- **Katacoda**
  - [Docker](#)
  - [Kubernetes](#)



# MORE FROM CARLOS

- **24 Hours of PASS**
  - [YouTube recording](#)
- **Simple Talk**
  - [SQL Server Docker Containers in macOS](#)
- **SQL Server Central**
  - [Creating Aliases for Docker commands](#)
  - [Managing SQL Server containers using Python – Part 1](#)
  - Managing SQL Server containers using Python – Part 2 (**Coming soon**)
- **MSSQL Tips**
  - [SQL Server 2019 CT2 RHEL Docker containers](#)





croblesdba



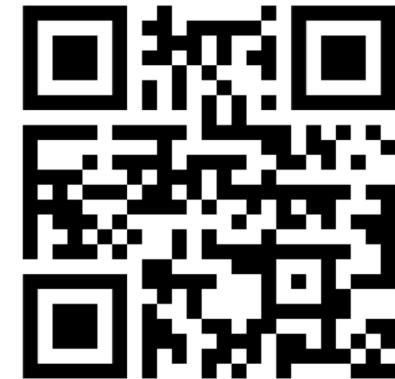
@dbamastery



crobles@dbamastery.com



DBA Mastery



# Thank you!!

